

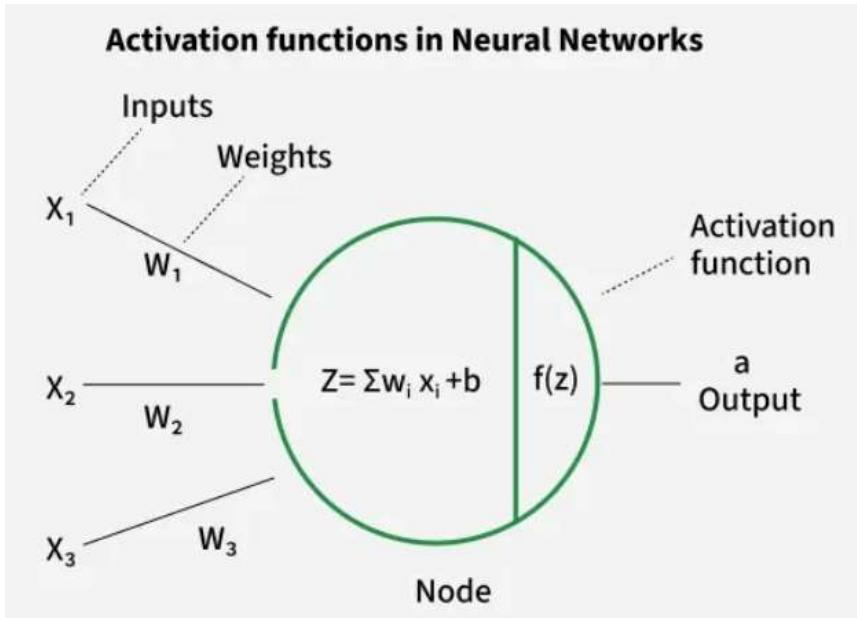
# Types of Activation Functions

CT-466 | Week 1 - Lecture 2

Instructor: Mehar Fatima Shaikh

# Activation function

- ▶ An activation function decides whether a neuron should be "activated" or not by applying a transformation to the weighted sum of inputs.
- ▶ It introduces non-linearity into the network, allowing deep learning models to learn complex patterns instead of just simple linear relationships.



# Why do we need Activation Functions?

- ▶ Without them, a neural network would just be a linear model, no matter how many layers it has.
- ▶ With activation functions, neural networks can model complex mappings (e.g., speech recognition, image classification).
- ▶ They help with gradient flow during backpropagation.

# Types of Activation Functions

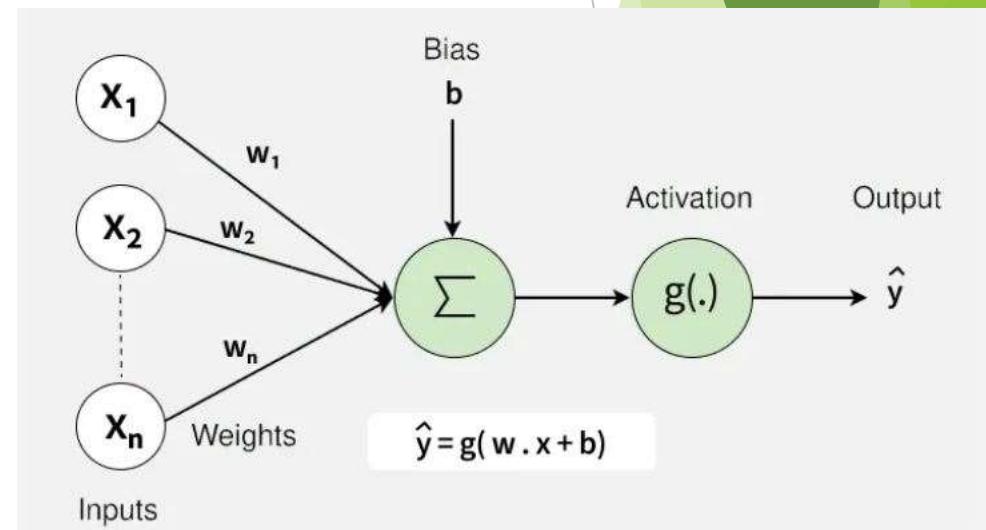
- ▶ Sigmoid Activation Function
- ▶ Tanh Activation Functions
- ▶ ReLU and Leaky ReLU Activation Functions

# Sigmoid activation function

- ▶ The **sigmoid activation function** is a mathematical function often used in neural networks.
- ▶ It looks like this:

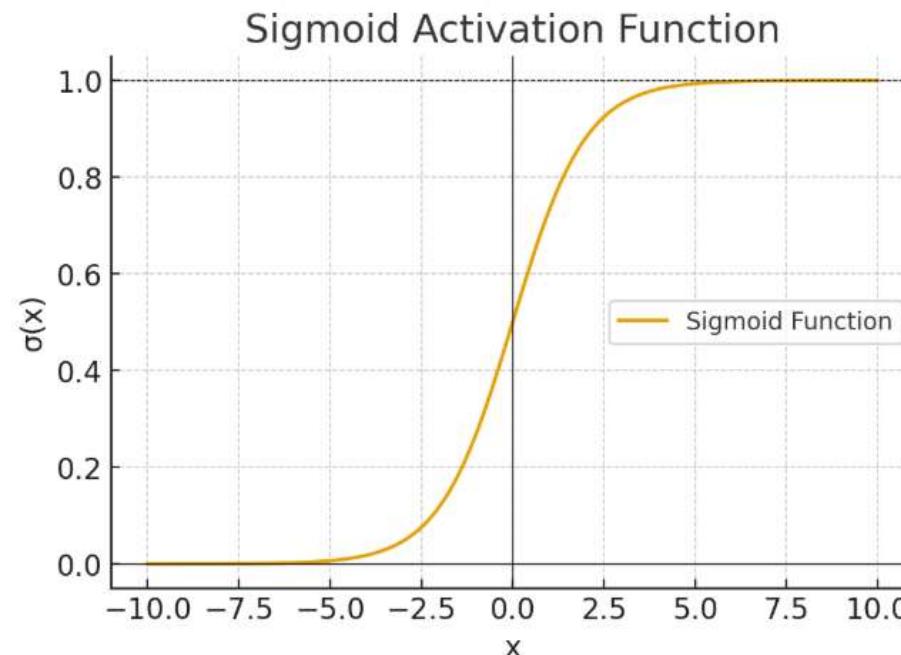
$$f(y) = \sigma(y) = \frac{1}{1 + e^{-y}}$$

- ▶ Applies on output layer.
- ▶ Range: Output is always between 0 and 1.
- ▶ Used for binary output
- ▶ Small inputs (negative x) → output is close to 0.
- ▶ Large inputs (positive x) → output is close to 1.
- ▶ At  $x = 0$  → output is 0.5.
- ▶ Since output is always Positive it creates bias.



# Why it's used:

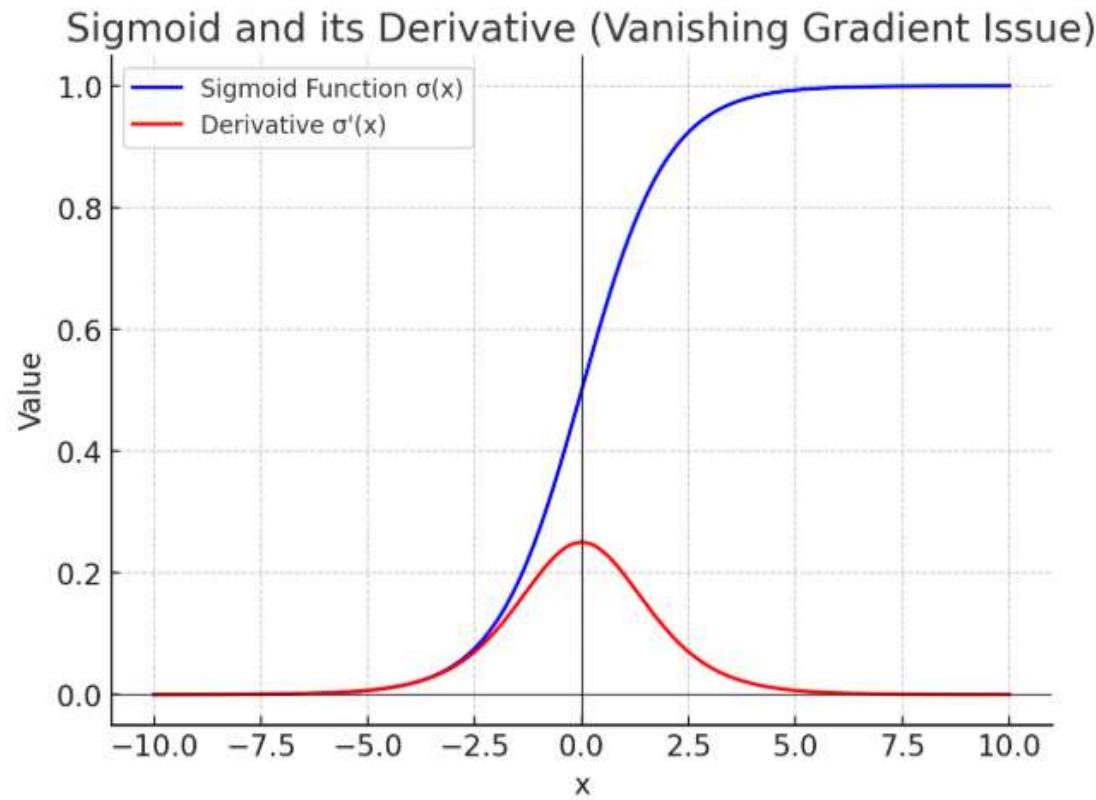
- ▶ Converts any input into a probability-like value between 0 and 1.
- ▶ Useful for binary classification problems.
- ▶ Limitation: It can cause vanishing gradients in deep networks, so nowadays ReLU and other functions are often preferred.



# Sigmoid Derivation

- ▶  $f(y) = \sigma(y) = \frac{1}{1+e^{-y}}$
- ▶  $\sigma'(x) = \frac{d f(y)}{dy} = \frac{d}{dy} \left( \frac{1}{1+e^{-y}} \right) = (1 + e^{-y})^{-1}$
- ▶  $\sigma'(x) = (-1) (1 + e^{-y})^{-2} (-e^{-y})$
- ▶  $\sigma'(x) = \frac{e^{-y}}{(1+e^{-y})^2} = \frac{e^{-y}}{1+e^{-y}} * \frac{1}{1+e^{-y}}$
- ▶  $\sigma'(x) = \frac{1+e^{-y}-1}{1+e^{-y}} * \frac{1}{1+e^{-y}} = \left( \frac{1+e^{-y}}{1+e^{-y}} - \frac{1}{1+e^{-y}} \right) * \frac{1}{1+e^{-y}}$
- ▶  $\sigma'(x) = (1 - \sigma(y)) * \sigma(y)$

# Sigmoid derivative with the vanishing gradient problem.



# Sigmoid derivative with the vanishing gradient problem.

$$\sigma'(x) = \sigma(x) \cdot (1 - \sigma(x))$$

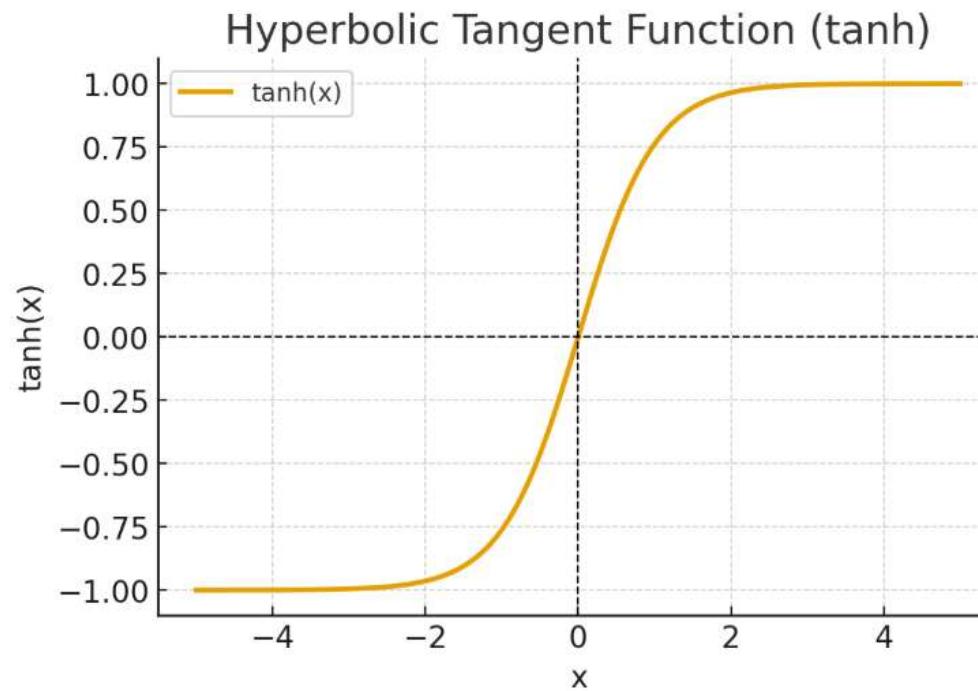
- ▶ Maximum value of derivative is 0.25 (when  $x=0$ , since  $\sigma(0)=0.5$ ).
- ▶ For large positive or negative  $x$ , the derivative becomes very close to 0.
- ▶ In deep neural networks, backpropagation multiplies many derivatives together. If each derivative is small (like 0.1 or 0.01), then after multiplying across many layers:

$$0.1^{10} = 1e-10 (\text{almost } 0!)$$

- ▶ So, the gradients shrink towards 0, and earlier layers stop learning → this is the vanishing gradient problem.

# Tanh (hyperbolic tangent function)

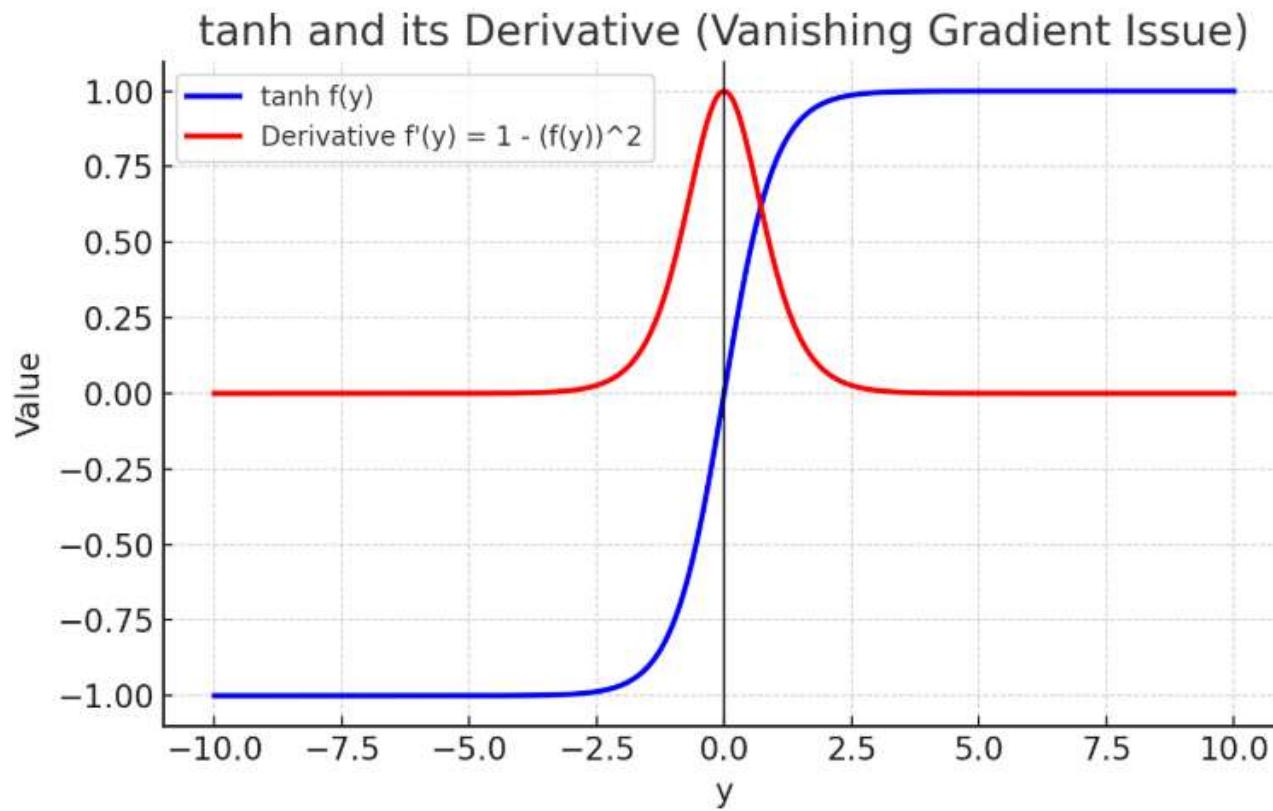
- ▶ In machine learning and neural networks, tanh is commonly used as an activation function because it outputs values in the range  $[-1,1]$ , making it useful for handling both positive and negative inputs.
- ▶ It is a mathematical function often written as:
- ▶  $f(y) = \tanh(y) = \frac{\sinh(y)}{\cosh(y)} = \frac{e^y - e^{-y}}{e^y + e^{-y}}$



# Derivative of Tanh

- ▶ Let  $f(y) = \tanh(y) = \frac{(e^y - e^{-y})}{(e^y + e^{-y})}$
- ▶ Define  $u(y) = (e^y - e^{-y})$  and  $v(y) = (e^y + e^{-y})$
- ▶ Then  $f(y) = \frac{u(y)}{v(y)}$
- ▶  $U'(y) = (e^y - (-1)e^{-y}) = (e^y + e^{-y}) = v(y)$
- ▶  $V'(y) = (e^y + e^{-y}) = (e^y + (-1)e^{-y}) = (e^y - e^{-y}) = u(y)$
- ▶  $f'(y) = \frac{(u'(y)v(y) - u(y)v'(y))}{v(y)^2} = \frac{(v(y)*v(y) - u(y)*u(y))}{v(y)^2} = \frac{(v(y))^2 - (u(y))^2}{v(y)^2}$
- ▶  $f'(y) = \frac{(v(y))^2 - (u(y))^2}{v(y)^2} = \frac{(e^y + e^{-y})^2 - (e^y - e^{-y})^2}{(e^y + e^{-y})^2}$
- ▶  $f'(y) = \frac{(e^y + e^{-y})^2}{(e^y + e^{-y})^2} - \frac{(e^y - e^{-y})^2}{(e^y + e^{-y})^2}$
- ▶  $f'(y) = 1 - (\frac{e^y - e^{-y}}{e^y + e^{-y}})^2$
- ▶  $f'(y) = 1 - (f(y))^2$

# Tanh derivative with the vanishing gradient problem.



# Vanishing gradient problem.

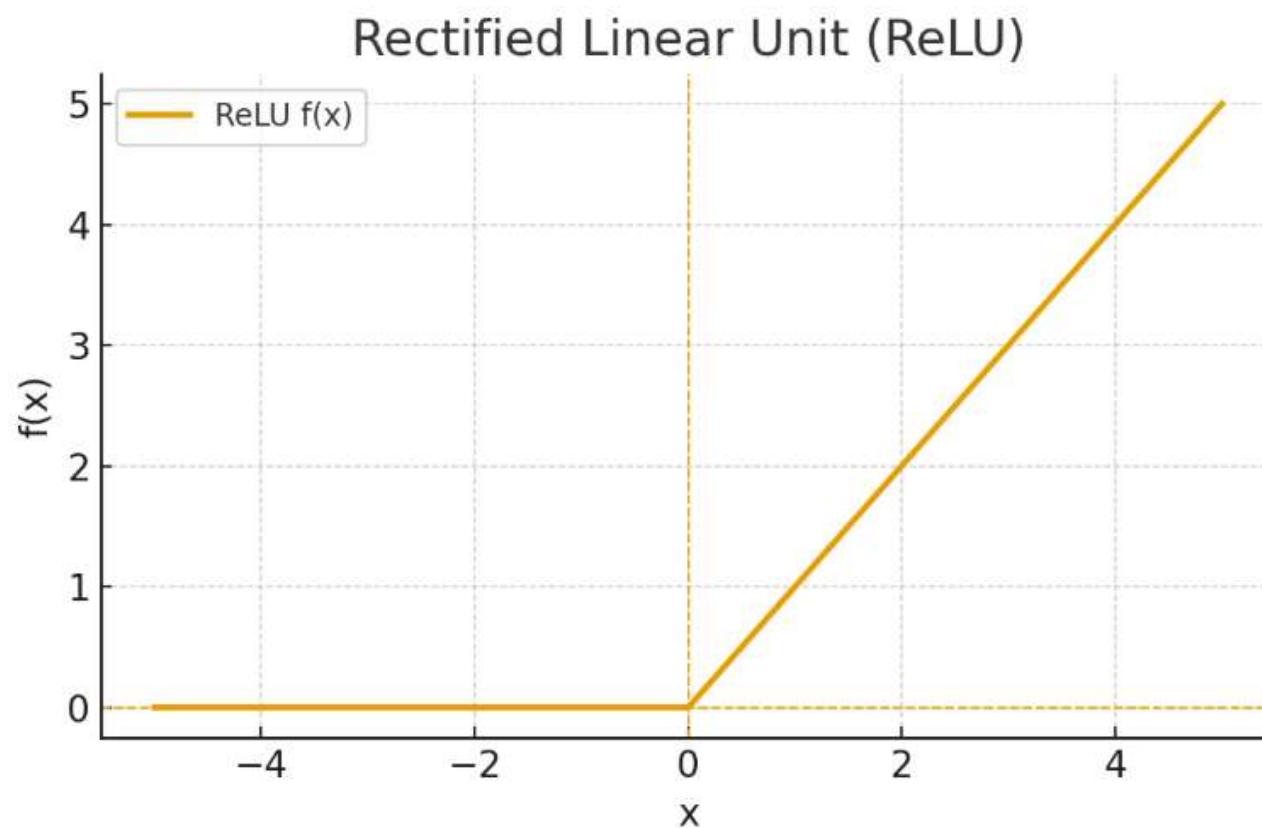
- ▶ The tanh function takes any real input and squashes it smoothly between  $-1$  and  $+1$ .
- ▶ Its derivative (slope) is largest when the input is near  $0$ . At that point, the slope can be close to  $1$ , so gradients flow well during backpropagation.
- ▶ But as the input moves far away (very positive or very negative), the tanh curve flattens. The slope (derivative) then becomes almost  $0$ .
- ▶ When training a deep neural network, gradients are multiplied layer after layer during backpropagation. If each derivative is a small number (close to  $0$ ), the product becomes tiny. This is called the vanishing gradient problem.
  - ▶ It means that early layers stop learning because their weight updates shrink almost to zero.
  - ▶ In practice, this slows down or completely blocks training of deep networks.
- ▶ tanh helps center data between  $-1$  and  $+1$ , but still suffers from vanishing gradients when inputs are far from  $0$ .

# ReLU (Rectified Linear Unit)

- ▶ The **ReLU (Rectified Linear Unit) activation function** is one of the most widely used functions in deep learning.
- ▶ It looks like this:

- ▶  $f(y)=\max(0,y)$

- ▶ 
$$f(x) = \begin{cases} x, & x > 0 \\ 0, & x \leq 0 \end{cases}$$

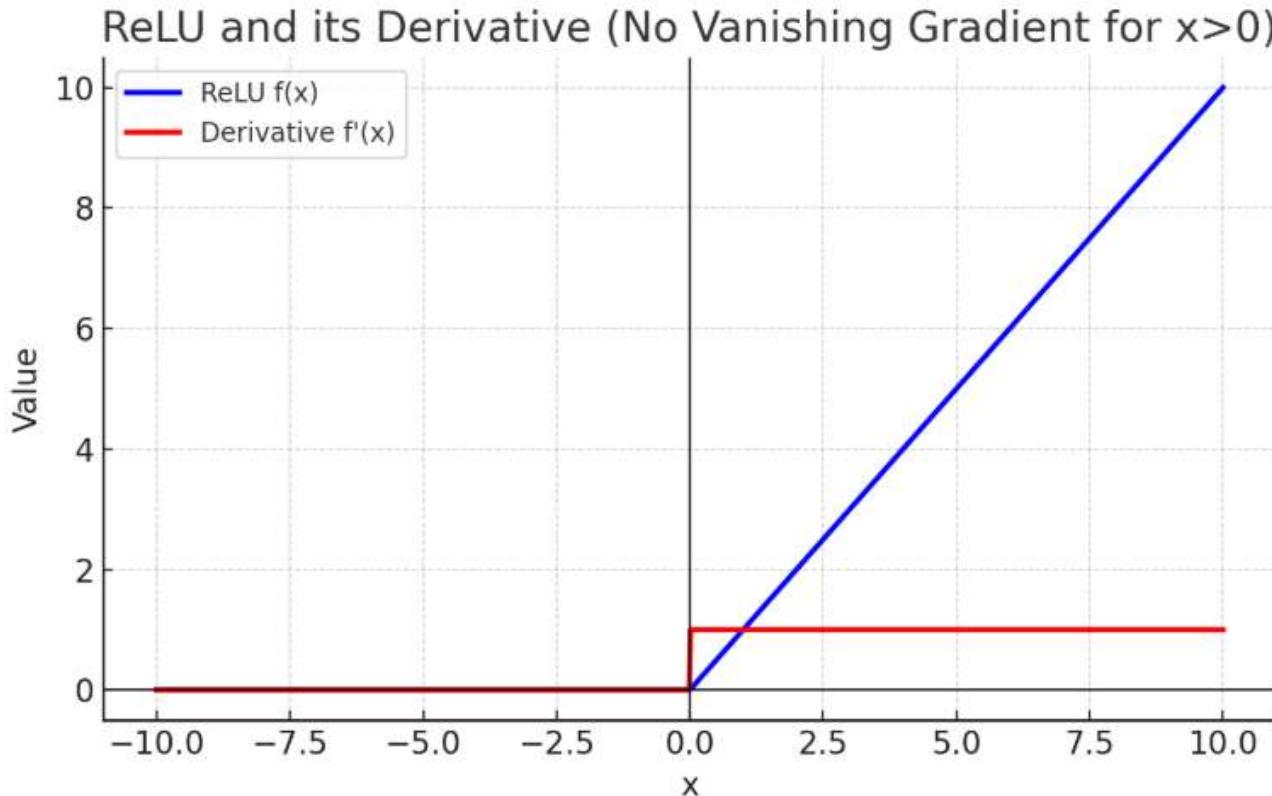


# ReLU (Rectified Linear Unit)

- ▶ Applies on hidden layers (commonly used in CNNs, ANNs, RNNs).
- ▶ Range: Output is between 0 and  $\infty$
- ▶ Used for introducing non-linearity while keeping computation simple.
- ▶ Small/negative inputs ( $y < 0$ )  $\rightarrow$  output is 0.
- ▶ Large/positive inputs ( $y > 0$ )  $\rightarrow$  output is equal to the input itself.
- ▶ At  $y=0 \rightarrow$  output is 0.
- ▶ Does not saturate for positive inputs, so it helps avoid vanishing gradients.
- ▶ Can cause the “dying ReLU” problem, where many neurons output 0 and stop learning if inputs are negative.

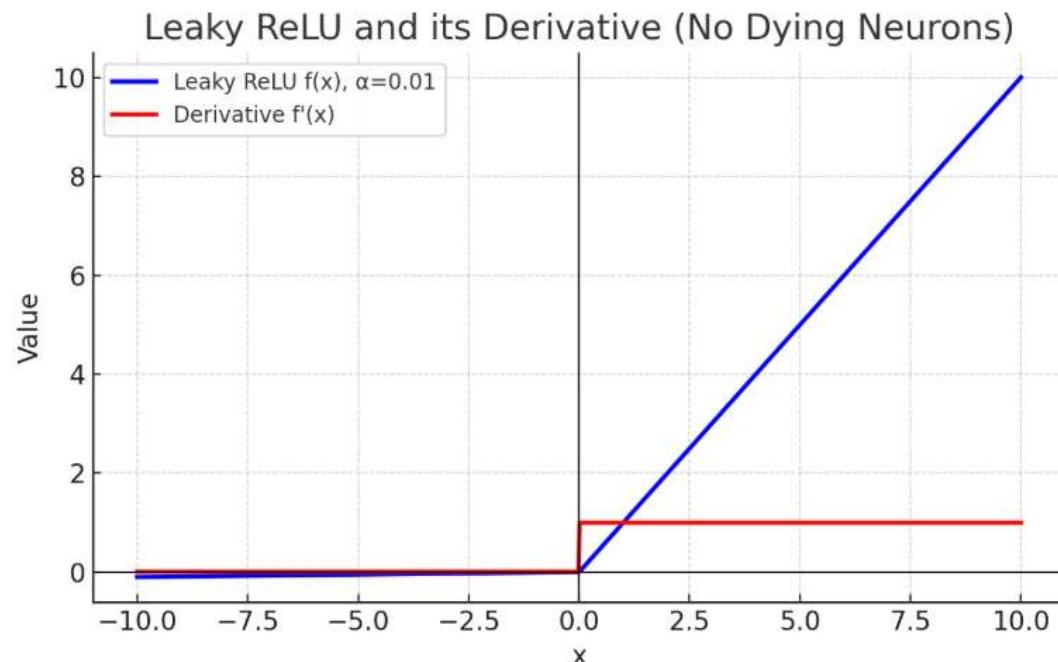
# ReLU derivative with the vanishing gradient problem.

- ▶  $f(x) = \begin{cases} 1, & x > 0 \\ 0, & x \leq 0 \end{cases}$



# Leaky ReLU

- ▶  $f(x) = \begin{cases} x, & x > 0 \\ \alpha x, & x \leq 0 \end{cases}$
- ▶ Where  $\alpha$  is a small constant (e.g., 0.01)
- ▶ For  $x > 0$ : behaves like ReLU (linear, slope = 1).
- ▶ For  $x \leq 0$ : instead of being flat, the function has a small negative slope ( $\alpha=0.01$ ), so the derivative is also small but nonzero.



# Python code

```
▶ # Inputs and weights
▶ x_inputs = [0.1, 0.5, 0.2]
▶ W_weights = [0.4, 0.3, 0.4]
▶ threshold = 0.5

▶ # Step function
▶ def step(weighted_sum):
▶     if weighted_sum > threshold:
▶         return 1
▶     else:
▶         return 0

▶ # Perceptron function
▶ def perceptron(x_inputs, W_weights):
▶     weighted_sum = 0
▶     for x, w in zip(x_inputs, W_weights):
▶         weighted_sum += x * w
▶     print("Weighted sum:", weighted_sum)
▶     return step(weighted_sum)

▶ # Run perceptron
▶ output = perceptron(x_inputs, W_weights)
▶ print("Output:", output)
```

```
1 # Inputs and weights
2 x_inputs = [0.1, 0.5, 0.2]
3 W_weights = [0.4, 0.3, 0.4]
4 threshold = 0.5
5
6 # Step function
7 def step(weighted_sum):
8     if weighted_sum > threshold:
9         return 1
10    else:
11        return 0
12
13 # Perceptron function
14 def perceptron(x_inputs, W_weights):
15     weighted_sum = 0
16     for x, w in zip(x_inputs, W_weights):
17         weighted_sum += x * w
18     print("Weighted sum:", weighted_sum)
19     return step(weighted_sum)
20
21 # Run perceptron
22 output = perceptron(x_inputs, W_weights)
23 print("Output:", output)
24 |
```

```
In [1]: runfile('C:/Users/N TECH/.spyder-py3/temp.py', wdir='C:/Users/N TECH/.spyder-py3')
Weighted sum: 0.27
Output: 0
```

```
In [2]:
```