

Non-linearly separable problems,

CT-466 | Week 5 - Lecture 10

Instructor: Mehar Fatima Shaikh

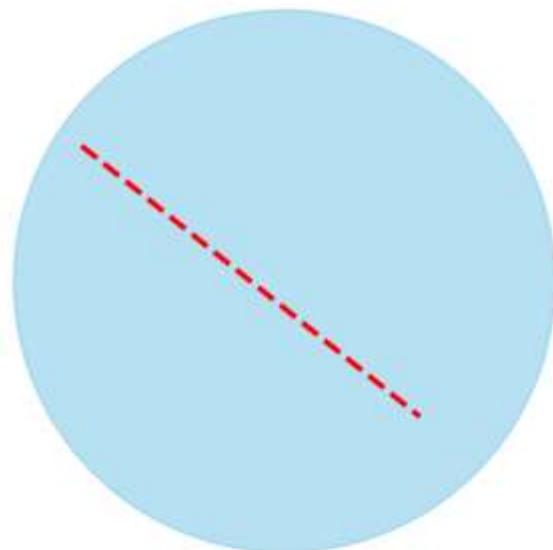
Convex Set

- ▶ A set S in a vector space is called convex if, for any two points $x_1, x_2 \in S$, the line segment joining them also lies entirely within the set.
- ▶ Mathematically:
 $\lambda x_1 + (1 - \lambda)x_2 \in S$ for all $x_1, x_2 \in S$ and $\lambda \in [0,1]$.

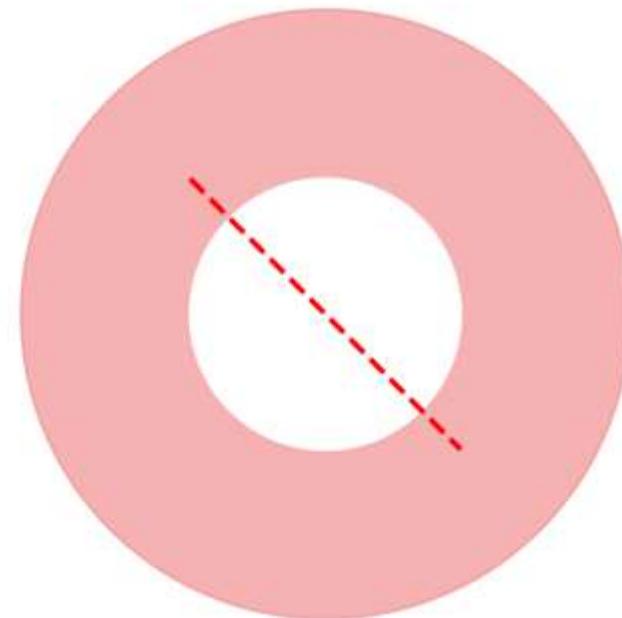
This means the set has no 'dents' or 'holes'; it's 'bulged outwards.'

Convex Set

Convex Set (Ellipse)



Non-Convex Set (Donut)



Convex Sets in ANN Context

In neural networks, convexity appears in several places:

1. Feasible Region of Parameters (Weights/Biases)

- During training, weights and biases form a parameter space.
- If we impose linear constraints (e.g., non-negative weights), the feasible solutions form a convex set.
- Example: If $w_1, w_2 \geq 0$, all valid weight pairs lie in the first quadrant, which is convex.

2. Loss Functions and Optimization

- For simple models like linear regression, the Mean Squared Error (MSE) loss is convex.
- Convexity of the loss function means no local minima → gradient descent will always find the global minimum.
- But in deep neural networks, due to multiple layers and nonlinear activation functions, the loss surface is generally non-convex (many local minima, saddle points).

3. Convex Sets in Regularization

- In L1 (Lasso) and L2 (Ridge) regularization, the constraint regions are convex:
 - L1 → diamond-shaped convex set
 - L2 → circular/elliptical convex set
- Convexity ensures optimization is tractable.

Why Convexity Matters in ANN

- ▶ Optimization simplicity: Convex sets + convex loss \rightarrow guarantees convergence to a unique solution.
- ▶ Deep networks: Loss surface is non-convex, but still contains large convex regions where gradient descent works well.
- ▶ Regularization: Convex constraint sets ensure stability and good generalization.

Convex Hull

- ▶ The **convex hull** of a set of points is the *smallest convex set* that contains all those points.
- ▶ Geometrically, imagine stretching a rubber band around the outermost points in a scatter plot → when released, it forms the convex hull.
- ▶ Mathematically:
 - ▶ $\text{Convex Hull}(X) = \{ \sum_{i=1}^n \lambda_i x_i \mid x_i \in X, \lambda_i \geq 0, \sum \lambda_i = 1 \}$

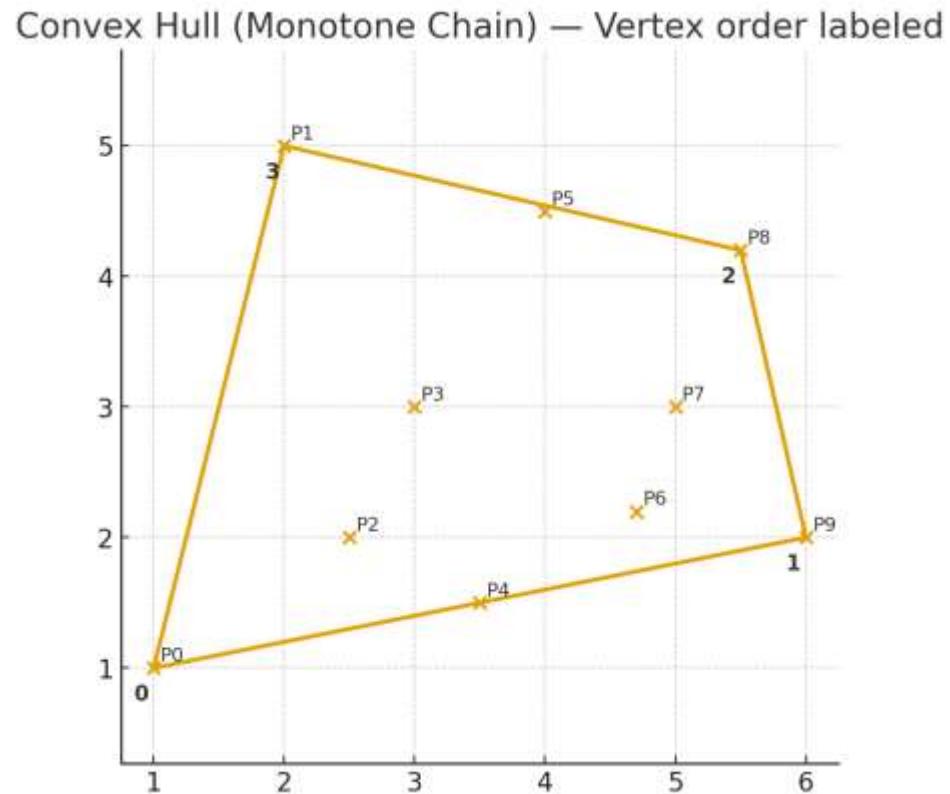
1) Pencil-and-paper (Gift Wrapping / Jarvis March)

- ▶ Goal: Walk around the “outer fence” of the points.
- ▶ 1. Plot points on graph paper (or a coordinate plane).
- ▶ 2. Start point: Pick the left-most point (if tie, pick the lowest y). Call it p0.
- ▶ 3. Pick direction: Imagine a horizontal ray to the left from p0 (or any fixed reference direction).
- ▶ 4. Find next vertex: For every other point q, compute the orientation of triangle (current, candidate, q). Use cross product sign:
$$\text{cross}(a, b, c) = (b_x - a_x)(c_y - a_y) - (b_y - a_y)(c_x - a_x)$$
Positive → left turn (more counter-clockwise). Choose the most counter-clockwise point as the next hull vertex.
- ▶ 5. Draw edge to that point and set it as current.
- ▶ 6. Repeat step 4 until you return to the starting point p0.
- ▶ 7. The edges you traced form the convex hull.
- ▶ Tip: If two points are collinear in the chosen direction, keep the farther one for the hull.

2) Algorithmic (Andrew's Monotone Chain)

- ▶ **Idea:** Sort points, then build **lower** and **upper** hulls using a left-turn test.
- ▶ **Steps**
- ▶ **Sort** points by x , then y .
- ▶ **Lower hull:**
 - ▶ Scan left→right.
 - ▶ While the last two points in the hull plus the new point make a **non-left turn** (cross ≤ 0), **pop** the last point.
 - ▶ Append the new point.
- ▶ **Upper hull:**
 - ▶ Scan right→left; do the same procedure.
- ▶ **Concatenate** lower[:-1] + upper[:-1] to avoid duplicating endpoints.
- ▶ The result is the convex hull in order (counter-clockwise).
- ▶ **Orientation test (same as above):**
- ▶ $\text{cross}(o \cdot a \cdot b) = (a_x - o_x)(b_y - o_y) - (a_y - o_y)(b_x - o_x)$
- ▶ Keep points that make **left turns** (cross > 0).

Andrew's Monotone Chain



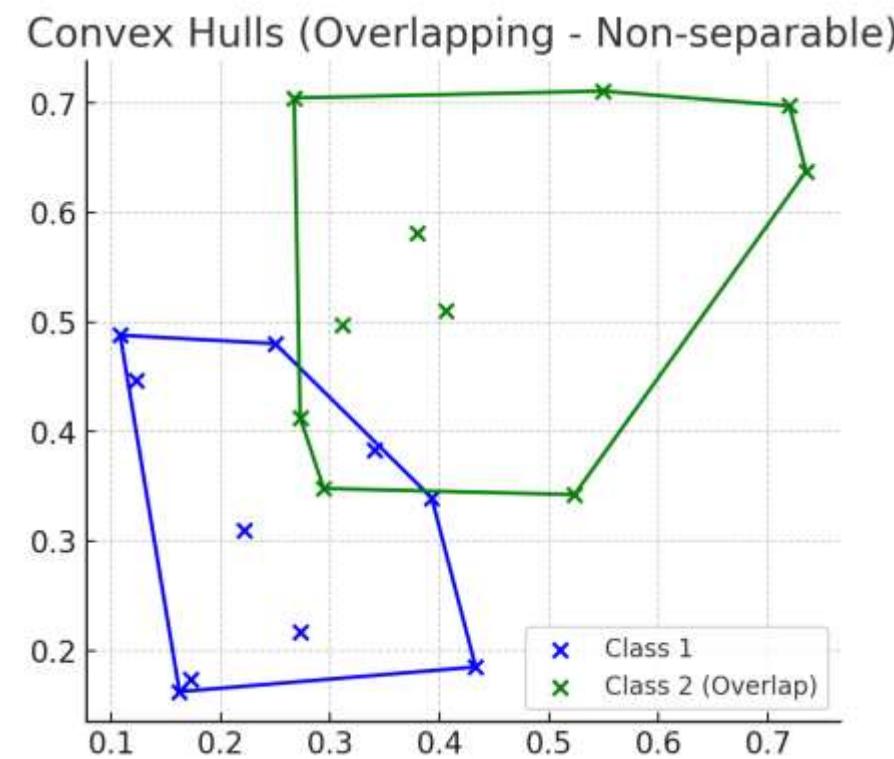
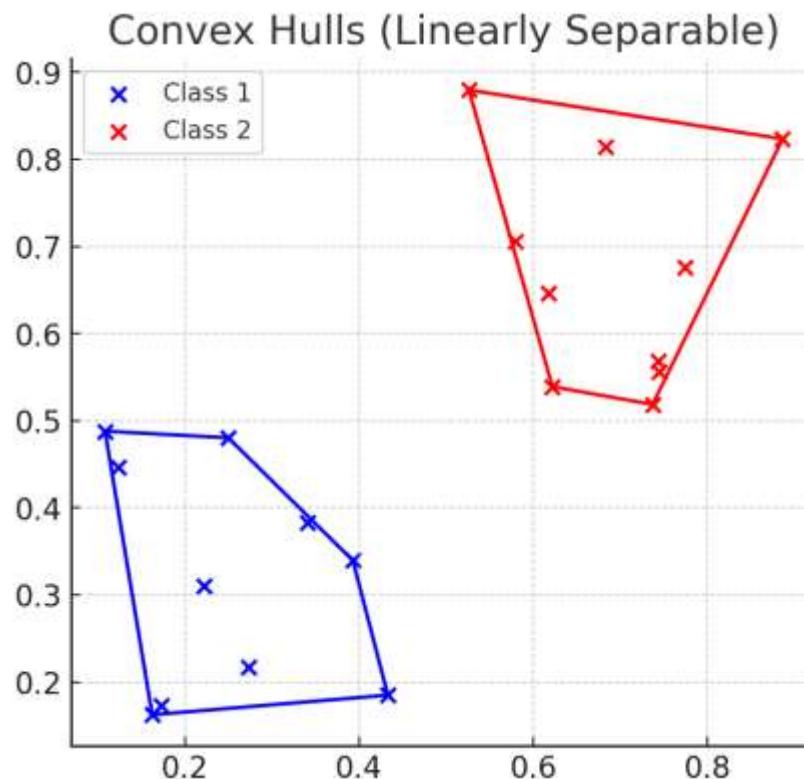
Convex Hull in ANN Context

- ▶ **Classification & Decision Boundaries**
 - ▶ In simple perceptron learning, if two classes' convex hulls **do not overlap**, the data is *linearly separable*.
 - ▶ That's why perceptrons can solve AND/OR problems (convex hulls separated), but **fail on XOR** (convex hulls overlap).
- ▶ **Geometry of Feature Space**
 - ▶ Convex hulls help visualize how training data is distributed in high-dimensional feature spaces.
 - ▶ ANN decision boundaries often approximate or cut through convex hulls.
- ▶ **Support Vector Machines (SVMs) relation**
 - ▶ SVM margin is effectively defined by support vectors lying on the convex hulls of different classes.

Why Convex Hull Matters in Neural Networks

- ▶ **Linearly separable data:** Perceptron convergence theorem says a perceptron will find a separating hyperplane *if convex hulls of the classes do not intersect.*
- ▶ **Non-linearly separable data** (like XOR): Convex hulls overlap → need Multi-Layer Perceptron's (MLPs) or kernel tricks.

Convex Hull



Linearly separable problems

- ▶ A problem is **linearly separable** if there exists a straight line (in 2D), a plane (in 3D), or a hyperplane (in higher dimensions) that can separate the data points of different classes **without error**.
- ▶ Mathematically:
- ▶ Dataset with inputs $x \in \mathbb{R}^n$ and labels $y \in \{-1, +1\}$
- ▶ Linearly separable if \exists weight vector w and bias b such that:
- ▶ $y_i(w \cdot x_i + b) > 0 \forall i$

Examples

- ▶ **Linearly Separable (can be solved by Perceptron):**
 - ▶ **AND Gate**

Inputs: $\{(0,0), (0,1), (1,0), (1,1)\}$
Output: Only $(1,1) \rightarrow 1$, others $\rightarrow 0$
 \checkmark Can be separated by a line.
 - ▶ **OR Gate**

Output is 0 only when $(0,0)$.
 \checkmark Can be separated by a line.
- ▶ **Not Linearly Separable (cannot be solved by single perceptron):**
- ▶ **XOR Gate**

$(0,0)$ and $(1,1)$ belong to one class, $(0,1)$ and $(1,0)$ belong to another.
No single straight line can separate these two groups.
Requires **Multi-Layer Perceptron (MLP)**.

Perceptron Convergence Theorem (PCT)

- ▶ Theorem Statement
- ▶ The Perceptron Convergence Theorem states:
- ▶ If a dataset is linearly separable, then the perceptron learning algorithm is guaranteed to find a separating hyperplane in a finite number of steps.
- ▶ Given a dataset with feature vectors $x_i \in \mathbb{R}^n$ and labels $y_i \in \{-1, 1\}$:
- ▶ If \exists weight vector w^* and bias b^* such that:
- ▶ $y_i(w^* \cdot x_i + b^*) > 0 \forall i$
- ▶ → The perceptron learning algorithm will converge to a solution vector w after a finite number of weight updates.

Weight Update Rule

- ▶ $w^{(t+1)} = w^{(t)} + \eta y_i x_i$
- ▶ where:
- ▶ η = learning rate
- ▶ y_i =true label
- ▶ x_i =input vector
- ▶ This happens **only when the perceptron misclassifies** the input.

Implications

- ▶ **Guarantee of Convergence**
 - ▶ If the data is linearly separable → perceptron will always converge.
 - ▶ If not linearly separable → algorithm never converges (keeps updating weights).
- ▶ **No Guarantee of Optimal Margin**
 - ▶ The perceptron finds *a* separating hyperplane, but not necessarily the one with **maximum margin** (unlike SVMs).
 - ▶ Each update “nudges” the hyperplane closer to correct classification.
 - ▶ Since the dataset is linearly separable, repeated nudging eventually aligns the hyperplane to separate the classes.

Example

- ▶ **AND Gate** (linearly separable): PCT ensures perceptron will converge.
- ▶ **XOR Gate** (not linearly separable): PCT does **not apply**, perceptron won't converge → need MLP.
- ▶ The **Perceptron Convergence Theorem** gives a mathematical guarantee: *for linearly separable problems, a perceptron will always converge in finite steps.* It is one of the earliest proofs that machine learning models can *learn from data* successfully.

Non-linearly separable problems

- ▶ A problem is **non-linearly separable** if no single straight line (in 2D), plane (in 3D), or hyperplane (in higher dimensions) can perfectly separate the data points of different classes.
- ▶ Mathematically, if there does **not** exist a weight vector w and bias b such that:
- ▶ $y_i(w \cdot x_i + b) > 0 \forall i$
- ▶ then the problem is non-linearly separable.

Examples

► XOR Gate

- ▶ Inputs: (0,0), (0,1), (1,0), (1,1)
- ▶ Outputs: 0, 1, 1, 0
No straight line can separate the two classes.
Needs **Multi-Layer Perceptron (MLP)**.

► Circular Decision Boundary

- ▶ Example: Class A = points inside a circle, Class B = points outside.
Cannot be separated by a line; requires a curved boundary.

► Spirals or Moons Dataset

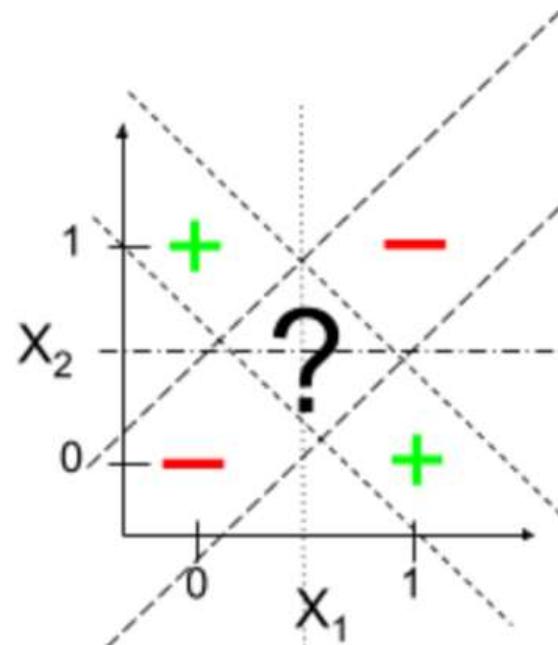
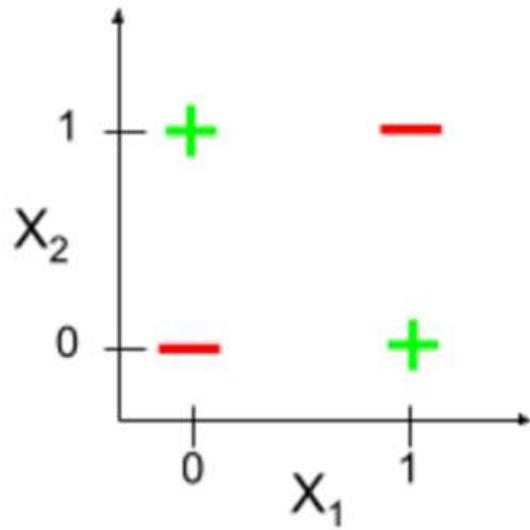
- ▶ Popular toy problems where data is distributed in non-linear shapes.
Require nonlinear models.

In ANN Context

- ▶ **Single Perceptron Limitation:**
A perceptron can only form a **linear decision boundary** → works only for linearly separable data.
- ▶ **Multi-Layer Perceptron (MLP):**
By stacking hidden layers + nonlinear activation functions (ReLU, sigmoid, tanh), MLPs can model **non-linear decision boundaries** and solve non-linearly separable problems.
- ▶ **Universal Approximation Theorem:**
An MLP with at least one hidden layer and nonlinear activations can approximate any continuous function → hence can handle non-linear problems.

XOR Problem

- ▶ A single perceptron draws one straight line (a linear decision boundary) in the (x_1, x_2) plane. For XOR, the 1's are at $(0,1)$ and $(1,0)$, while the 0's are at $(0,0)$ and $(1,1)$. No single line can separate those two classes hence non-linear separability.
(Geometric intuition: the positive points sit on opposite corners; any straight line that groups them together also includes at least one negative point.).



Why Single-Layer Perceptron Fails

- ▶ **Linear Decision Boundary**
 - ▶ A perceptron computes:
- ▶ $y = f(w_1x_1 + w_2x_2 + b)$
 - ▶ This is always a **linear separator** (straight line in 2D).
- ▶ **XOR is Non-Linearly Separable**
 - ▶ The two classes in XOR are **diagonal** to each other.
 - ▶ Their convex hulls overlap → no linear hyperplane exists that separates them.
 - ▶ Hence, perceptron keeps updating weights forever and never converges.

Limitations of Single-Layer Perceptron

- ▶ Can only solve linearly separable problems (like AND, OR, NOT).
- ▶ Cannot solve non-linear problems (like XOR, circle-in-circle, spirals).
- ▶ Provides only linear decision boundaries.
- ▶ Limited computational power → not a universal function approximator.

Solution: Multi-Layer Perceptron (MLP)

- ▶ Introduce hidden layers + nonlinear activation functions.
- ▶ An MLP can combine linear boundaries into nonlinear ones, making it capable of solving XOR and other complex problems.
- ▶ This was the historical reason why backpropagation and deep networks became essential in AI.